

Handed out: 11/15/2017 Due by 11:59 PM, midnight (CST) on Tuesday, 11/21/2017

You can do this assignment on your PC, Mac, Linux on Windows or Linux VM. Tell us which environment you are using. Implementation of this assignment will require some light reading of Azure literature on the Internet. If the assignment refers to CLI, you are welcome to use Power Shell or even one of programming languages if that is more convenient for you.

Problem 01. Create an Azure blob container and move the attached file `storage_table_demo.py` as a Blob into that container. Demonstrate that you can modify the file and upload it to the same container again. This time download the file back to your operating system and prove that you received back file with recent modifications. Delete the blob, container and the storage account. Do it all using Azure CLI.
(15%)

CREATE RESOURCE GROUP - RG-KIRK

```
kirks-mpb:~ el5vgxz$ az group create --name rg-kirk --location "West US"
{
```

CREATE STORAGE ACCOUNT (truncated output)

```
kirks-mpb:~ el5vgxz$ az storage account create --name sakirk -g rg-kirk -l "West US"
--sku Standard_LRS --encryption blob
{
  "accessTier": null,
  "creationTime": "2017-11-20T19:22:23.598668+00:00",
  "customDomain": null,
  "enableHttpsTrafficOnly": false,
  "encryption": {
    "keySource": "Microsoft.Storage",
    "keyVaultProperties": null,
    "services": {
      "blob": {
        "enabled": true,
        "lastEnabledTime": "2017-11-20T19:22:23.605669+00:00"
      },
      "file": {
        "enabled": true,
        "lastEnabledTime": "2017-11-20T19:22:23.605669+00:00"
      },
      "queue": null,
      "table": null
    }
  },
  "id": "/subscriptions/d5f65876-4dec-4a2e-afe1-b61e713a8612/resourceGroups/rg-kirk/providers/Microsoft.Storage/storageAccounts/sakirk",
```

DISPLAY KEYS

```
kirks-mbp:~ el5vgxz$ az storage account keys list --account-name sakirk --resource-group rg-kirk --output table
```

KeyName	Permissions	Value
key1	Full	9lsYtzs1kG0dSULPLUXgfh+V/A/rTVHa1wozF/tVyZzya1PsMNqAWddKxoBKzCv5h819Z3fcGk4bcWRbu7Aogg==
key2	Full	m2mBLXbHBD3/CAADwpow6WNKlUPLZp41hjDnXghXFqVw+7C0EMIArUfPgZh6/gvEsHoIH6YZezsLmvnCVNt5/w==

CREATE CONTAINER

```
kirks-mbp:~ el5vgxz$ az storage container create --name mystoragecontainer
{
  "created": true
}
```

UPLOAD FILE

```
kirks-mbp:Azure Training el5vgxz$ az storage blob upload --container-name mystoragecontainer --name myblob --file ./storage_table_demo.py
Finished[#####] 100.0000%
{
  "etag": "\"0x8D530509E10038C\"",
  "lastModified": "2017-11-20T19:55:15+00:00"
}
```

My blob service and container with file uploaded

The screenshot shows the Azure portal interface for a storage account named 'sakirk'. On the left, the 'Essentials' sidebar shows a list of containers, with 'mystoragecontainer' selected. The main pane displays the details for the 'mystoragecontainer' container. It shows the location as 'mystoragecontainer' and a search bar for blobs. Below the search bar, a table lists the blobs in the container:

NAME	MODIFIED	BLOB TYPE	SIZE	LEASE STATE
myblob	11/20/2017, 12:55:15 PM	Block blob	6.16 KIB	Available

MODIFY FILE ADDING LINE TO END – TAIL LAST 2 LINES – UPLOAD NEW FILE – MODIFICATION TIME ON FILE CHANGED

```
kirks-mbp:Azure Training el5vgxz$ echo "File modified by Kirk" >> storage_table_demo.py
kirks-mbp:Azure Training el5vgxz$ tail -n2 storage_table_demo.py
print('Error deleting resource group.')
File modified by Kirk
kirks-mbp:Azure Training el5vgxz$ az storage blob upload --container-name mystoragecontainer --name myblob --file ./storage_table_demo.py
Finished[#####] 100.0000%
{
  "etag": "\"0x8D530520B04AD87\"",
  "lastModified": "2017-11-20T20:05:27+00:00"
}
```

DELETE FILE AND DOWNLOAD FILE S “NEWFILE”

```
kirks-mbp:Azure Training el5vgxz$ rm storage_table_demo.py
kirks-mbp:Azure Training el5vgxz$ ls stora*
ls: stora*: No such file or directory
kirks-mbp:Azure Training el5vgxz$ az storage blob download --container-name mystoragecontainer --name myblob --file ./newfile
{
  "content": null,
  "metadata": {},
}
```

```
"name": "myblob",  
"properties": {
```

TAIL NEWFILE SHOWING MODIFICATIONS

```
kirks-mbp:Azure Training el5vgxz$ tail -n3 newfile  
else:  
    print('Error deleting resource group.')
```

File modified by Kirk

DELETE EVERYTHING INDIVIDUALLY

```
kirks-mbp:Azure Training el5vgxz$ az storage blob delete --container-name  
mystoragecontainer --name myblob  
{  
  "deleted": null  
}  
kirks-mbp:Azure Training el5vgxz$ az storage container delete --name  
mystoragecontainer  
{  
  "deleted": true  
}  
kirks-mbp:Azure Training el5vgxz$ az storage account delete --resource-group rg-kirk  
--name sakirk  
Are you sure you want to perform this operation? (y/n): y  
kirks-mbp:Azure Training el5vgxz$ az group delete --name rg-kirk  
Are you sure you want to perform this operation? (y/n): y  
kirks-mbp:Azure Training el5vgxz$
```

Problem 02. Examine attached file `storage_table_demo.py`. This is the code we discussed in class which creates and populates Azure Table structure called “itemstable”. Examine the code carefully. It contains practically all the tools you need to deal with Azure Tables. Replace prefix “zdz” of all object names with a short string unique for you. Rather than populate table with Pizzas, populate one of its partitions with cars, as if you are a car dealership. Cars are characterized by make, model, year, color and price. Populate yet another partition with coffee shop inventory. Coffee is characterized by the brand, flavor, size of the cup and price per cup. Place your modified file in your GitHub repository. Open Portal’s Cloud Shell and transfer file from your repository into your home directory of the Cloud Shell. Cloud Shell does have `git` installed. Run your Python program using command: `python2.7 storage_table_demo.py`. Capture the output. Notice that program waits for your commands. Before you delete your table `itemstable`, open Azure Storage Explorer and capture the content of your table. Could you think of another way of transferring files to your Cloud Shell home directory?
(20%)

FILE MODIFIED FOR CAR AND COFFEE

```
# A partition key tracks how like-minded entries in the Table are created and  
queried.  
# A row key is a unique ID for each entity in the partition  
# These two properties are used as a primary key to index the Table. This makes  
queries much quicker.  
  
car = Entity()  
car.PartitionKey = 'Inventory'  
car.RowKey = '001'
```

```

car.year = '2017'
car.make = 'Tesla'
car.model = 'Model S'
car.color = 'Black'
car.price = 72000
table_service.insert_entity('itemstable', car)
print('Created entry for Kirks car')

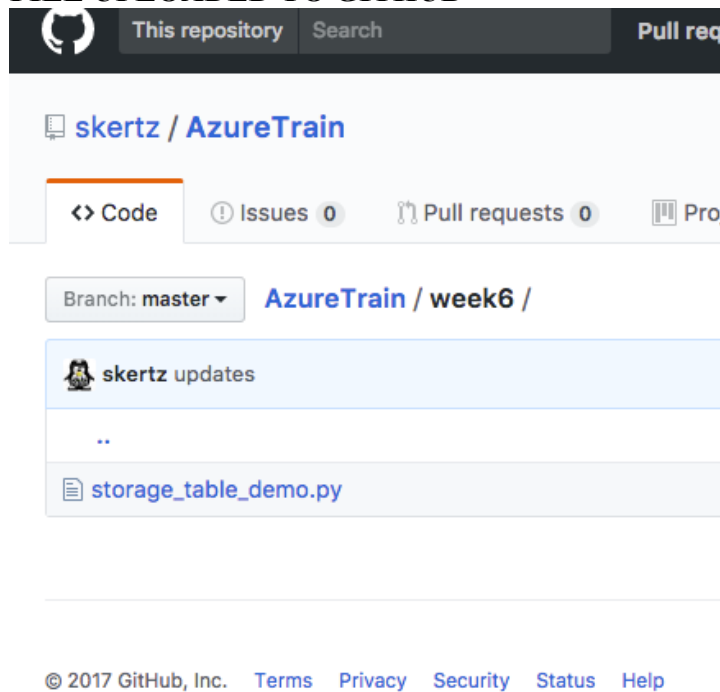
car.PartitionKey = 'Inventory'
car.RowKey = '002'
car.year = '2003'
car.make = 'BMW'
car.model = '330 ci'
car.color = 'Black'
car.price = 5000
table_service.insert_entity('itemstable', car)
print('Created entry for Matts car')

coffee = Entity()
coffee.PartitionKey = 'coffehouse'
coffee.RowKey = '002'
coffee.brand = 'Folders'
coffee.flavor = 'regular'
coffee.size = 'large'
coffee.cost = 1.99
table_service.insert_entity('itemstable', coffee)
print('Created entry for a large regualre Folders...\n')
time.sleep(1)

coffee = Entity()
coffee.PartitionKey = 'coffehouse'
coffee.RowKey = '001'
coffee.brand = 'Maxwell'
coffee.flavor = 'decaf'
coffee.size = 'small'
coffee.cost = .98
table_service.insert_entity('itemstable', coffee)
print('Created entry for a small decaf Maxwell...\n')
time.sleep(1)

```

FILE UPLOADED TO GITHUB



SETUP CLOUD SHELL AND GIT CLONE MY REPO

```
Bash
Type "az" to use Azure CLI 2.0
Type "help" to learn about Cloud Shell

kirk@Azure:~$ pwd
/home/kirk
kirk@Azure:~$ git clone https://github.com/skertz/AzureTrain.git
Cloning into 'AzureTrain'...
remote: Counting objects: 26, done.
remote: Compressing objects: 100% (18/18), done.
remote: Total 26 (delta 1), reused 23 (delta 1), pack-reused 3
Unpacking objects: 100% (26/26), done.
Checking connectivity... done.
kirk@Azure:~$ ls
AzureTrain  clouddrive
kirk@Azure:~$ cd AzureTrain/week6/
kirk@Azure:~/AzureTrain/week6$ ls
storage_table_demo.py
kirk@Azure:~/AzureTrain/week6$
```

EXECUTE PYTHON SCRIPT

```
kirk@Azure:~/AzureTrain/week6$ python2.7 storage_table_demo.py
Resource group: kedgrp9q1 created successfully.
```

Storage account: ked8lqwpj created successfully.

Let's create an Azure Storage Table to store some data.

Press Enter to continue...

Storage Table: itemstable created successfully.

Now let's add some entries to our Table.

Remember, Azure Storage Tables is a NoSQL datastore, so this is similar to adding records to a database.

Press Enter to continue...

Created entry for Kirks car

Created entry for Matts car

Created entry for Renees car

Created entry for a small regular Folders...

Created entry for a large regualre Folders...

Created entry for a small decaf Maxwell...

With some data in our Azure Storage Table, we can query the data.

Let's see what the car inventory looks like.

Press Enter to continue...

This is a basic example of how Azure Storage Tables behave like a database.

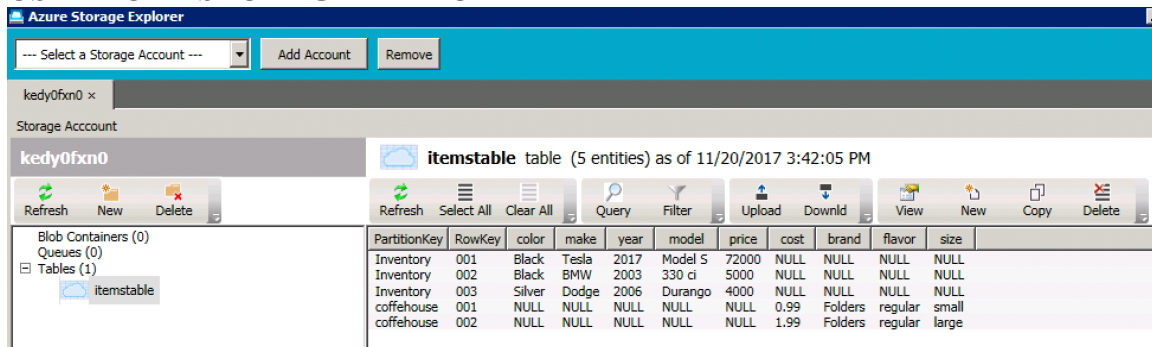
To keep things tidy, let's clean up the Azure Storage resources we created.

Press Enter to continue...

Storage table: itemstable deleted successfully.

Resource group: kedgrp9ql deleted successfully.

USE AZURE STORAGE EXPLORER



The screenshot shows the Azure Storage Explorer application. The top bar includes a dropdown for 'Select a Storage Account', 'Add Account', and 'Remove' buttons. Below this, the 'Storage Account' section shows 'kedy0fxn0'. The main pane displays the 'itemstable' table with 5 entities as of 11/20/2017 3:42:05 PM. The table has columns for PartitionKey, RowKey, color, make, year, model, price, cost, brand, flavor, and size. The data is as follows:

PartitionKey	RowKey	color	make	year	model	price	cost	brand	flavor	size
Inventory	001	Black	Tesla	2017	Model S	72000	NULL	NULL	NULL	NULL
Inventory	002	Black	BMW	2003	330 ci	5000	NULL	NULL	NULL	NULL
Inventory	003	Silver	Dodge	2006	Durango	4000	NULL	NULL	NULL	NULL
coffehouse	001	NULL	NULL	NULL	NULL	NULL	0.99	Folders	regular	small
coffehouse	002	NULL	NULL	NULL	NULL	NULL	1.99	Folders	regular	large

Problem 03. Using Azure CLI create a File Share. Move two or free files to that file share. Create two or three directories in that file share. Subsequently, mount that file share as a shared drive on your Windows or MacOS machine. Demonstrate that you can place files from your Windows or MacOS into one of the directories on the shared drive and that they will be visible in Azure Portal.

(20%)

CREATE RESOURCE GROUP

```
kirks-mbp:week6 el5vgxz$ az group create --name rg-kirk --location "West US"
```

CREATE STORAGE ACCOUNT (kind = file)

```
kirks-mbp:week6 el5vgxz$ az storage account create --kind Storage --resource-group rg-kirk --name sakirk --sku Standard_LRS {
  "accessTier": null,
  "creationTime": "2017-11-20T21:54:30.949268+00:00",
  "customDomain": null,
```

DISPLAY CONNECTION STRING

```
kirks-mbp:week6 el5vgxz$ az storage account show-connection-string -n sakirk -g rg-kirk --query 'connectionString' -o tsv
DefaultEndpointsProtocol=https;EndpointSuffix=core.windows.net;AccountName=sakirk;AccountKey=FbiiQ+9Lqkk06Kgc+qfbWBn0MEepiCIB6/uym7y+xWSQJm27Xt6IhDYqQI7wfuGCvE5yU9BvRI/UFE83ULQ4vQ==
```

CREATE FILE SHARE

```
kirks-mbp:week6 el5vgxz$ az storage share create --name kirkshare --account-name sakirk {
  "created": true
}
```

MOUNT SHARE (showing with df command)

```
kirks-mbp:week6 el5vgxz$ mount -t smbfs -o -f=0777,-d=0777 //sakirk2.file.core.windows.net/kirkshare azuremount

kirks-mbp:week6 el5vgxz$ df | grep azu
//sakirk2@sakirk2.file.core.windows.net/kirkshare 4294967296      128
4294967168      1%      2      134217724      0% /Users/el5vgxz/Desktop/code/AzureTrain/week6/azuremount
```

COPY 3 TEST FILES TO SHARE/MOUNT

```
kirks-mbp:week6 el5vgxz$ ls
azuremount      file1          file2          file3
storage_table_demo.py
kirks-mbp:week6 el5vgxz$ cp file* azuremount/
kirks-mbp:week6 el5vgxz$
```

SHOWING FILES EXIST ON SHARE

File Service sakirk2

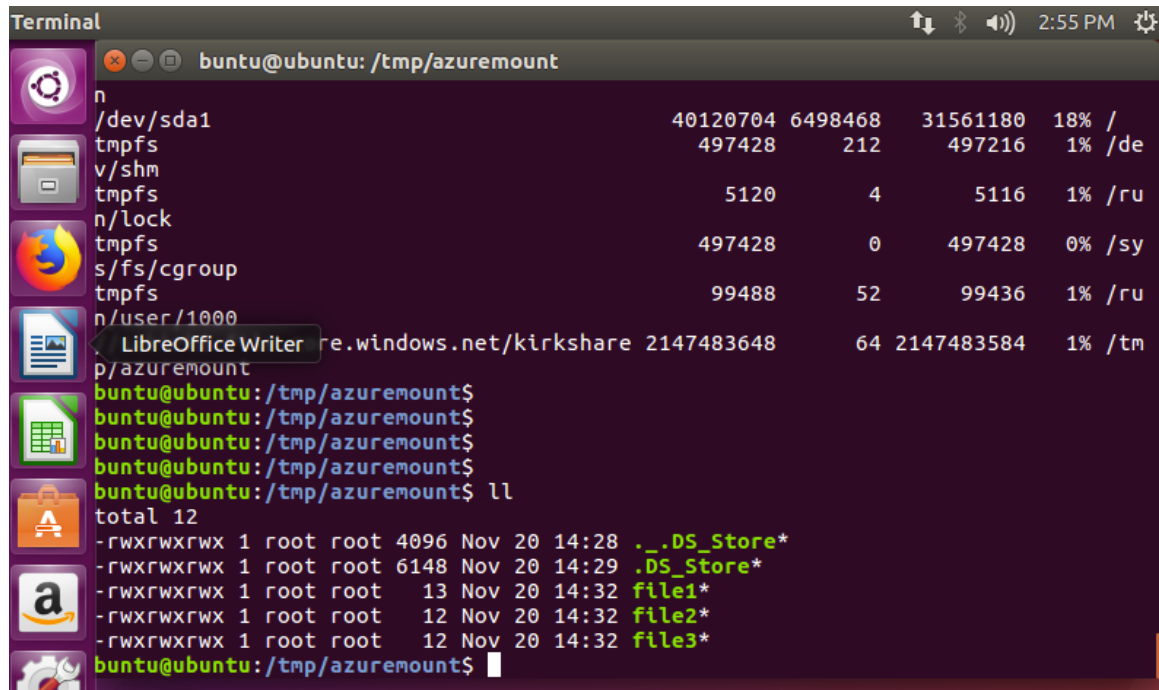
File share

Location: kirkshare

NAME	TYPE	SIZE
.DS_Store	File	6 KB
..DS_Store	File	4 KB
file1	File	13 B
file2	File	12 B
file3	File	12 B

Problem 04. Mount the file share you created in problem 3 as a shared drive on your Linux VM, either a CentOS or Ubuntu. This might require some Internet search. You need to find out how to open port 445 on your VM. Demonstrate that you can navigate to the shared folder, create a new directory using Linux `mkdir` command and copy a file from your Linux OS into that directory. Verify that the directory and the file are visible in Azure Portal. Could you find Azure CLI or Power Shell command to list the content of your Azure File Share
(20%)

MOUNTED SHARE IN UBUNTU – SEE DF COMMAND AND LS -L of existing files

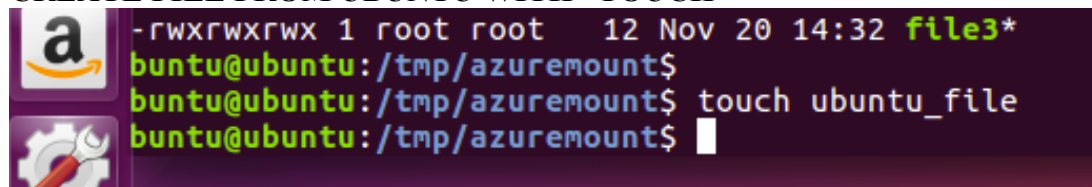


```

Terminal
buntu@ubuntu: /tmp/azuremount
df
Filesystem            Size  Used Avail Use% Mounted on
/dev/sda1              40G  4.1G   36G   10% /
tmpfs                  497M   212K  497M    1% /dev/shm
v/shm                  512K     4K   512K    1% /run/lock
tmpfs                  497M     0K  497M    0% /sys/fs/cgroup
tmpfs                  99488K   52K  99436K    1% /run/user/1000
LibreOffice Writer    re.windows.net/kirkshare 2147483648    64 2147483584    1% /tmp/azuremount
buntu@ubuntu: /tmp/azuremount$
buntu@ubuntu: /tmp/azuremount$
buntu@ubuntu: /tmp/azuremount$
buntu@ubuntu: /tmp/azuremount$
buntu@ubuntu: /tmp/azuremount$ ll
total 12
-rwxrwxrwx 1 root root 4096 Nov 20 14:28 ._.DS_Store*
-rwxrwxrwx 1 root root 6148 Nov 20 14:29 .DS_Store*
-rwxrwxrwx 1 root root  13 Nov 20 14:32 file1*
-rwxrwxrwx 1 root root  12 Nov 20 14:32 file2*
-rwxrwxrwx 1 root root  12 Nov 20 14:32 file3*
buntu@ubuntu: /tmp/azuremount$

```

CREATE FILE FROM UBUNTU WITH “TOUCH”

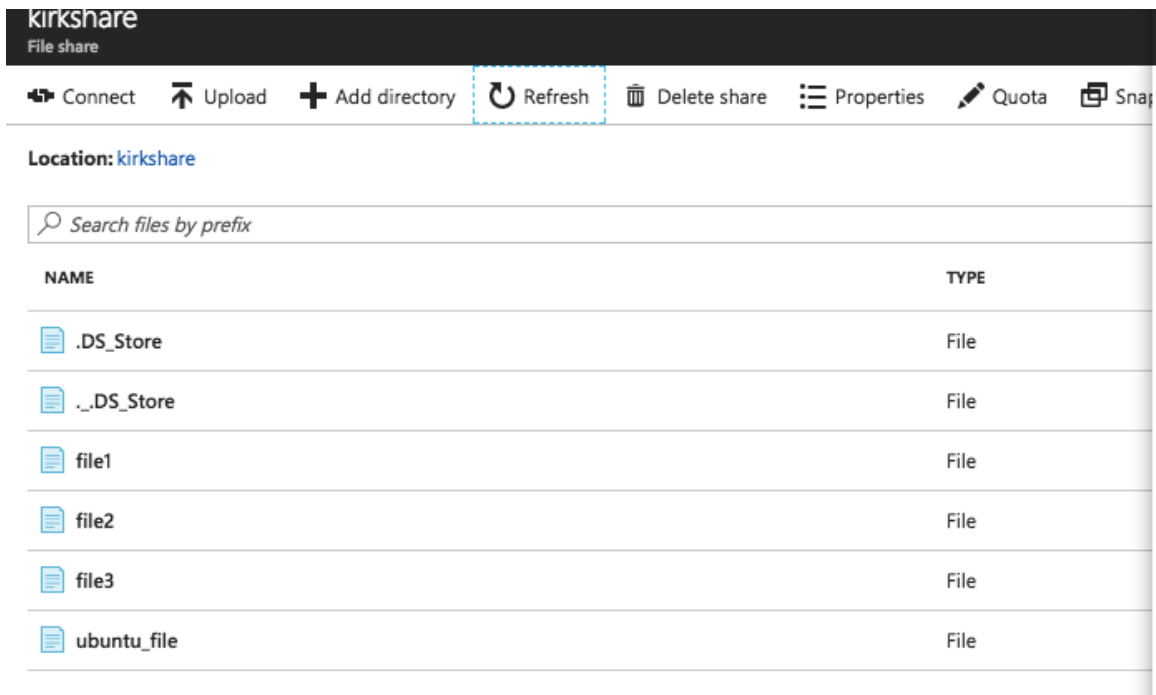


```

-rwxrwxrwx 1 root root  12 Nov 20 14:32 file3*
buntu@ubuntu: /tmp/azuremount$
buntu@ubuntu: /tmp/azuremount$ touch ubuntu_file
buntu@ubuntu: /tmp/azuremount$

```

FILE (ubuntu_file) IS AVAILABLE ON PORTAL



AZ CLI TO DISPLAY FILES IN SHARE

```
kirks-mbp:week6 el5vgxz$ az storage file list --share-name kirkshare --account-name sakirk2 | grep name
  "name": ".DS_Store",
  "name": "._DS_Store",
  "name": "file1",
  "name": "file2",
  "name": "file3",
  "name": "ubuntu_file",
```

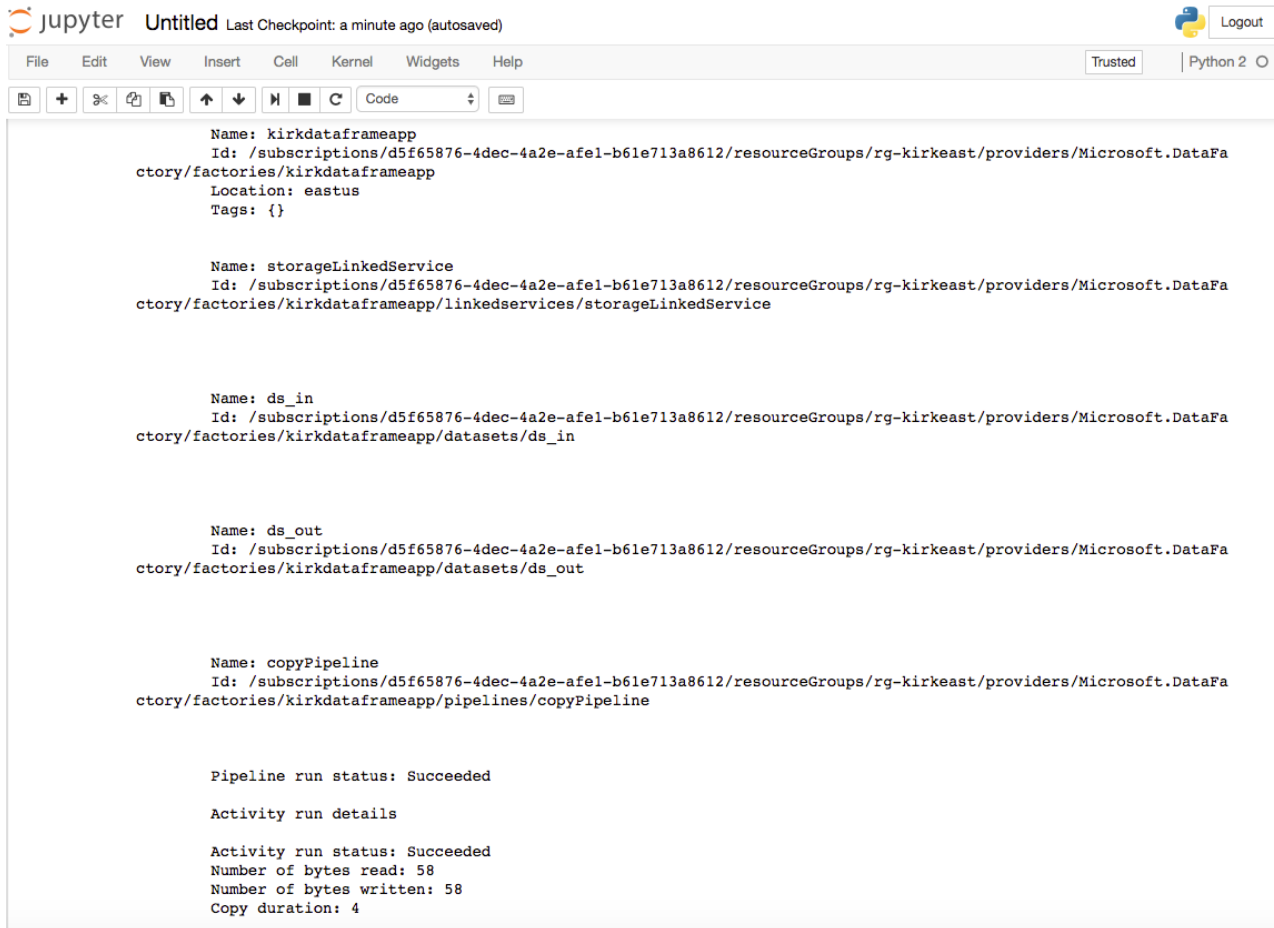
Problem 05. Follow the slides from Lab 06 on creation of a data factory for the Visual Studio (C#) or for Python. You are expected to follow one or the other approach. Create all resources noted in the slides. For the data create a file called `astroplayers.txt` with 5 players from the World Series Houston Astros Roster (first and last names) or use names from your favorite team. Create a container called `playerscontainer`. Upload `astroplayers.txt` to a folder in your container called `myteam`. After you run your code check that all of your 5 favorite players are in your SQL database table (Visual Studio) or your container's output folder (Python) using the Azure portal. (20%)

INPUT PARMS

AZ Subscription ID: d5f65876-4dec-4a2e-afe1-b61e713a8612
 Storage Name: saskertz234564
 Storage
 Key: mEw3N5/jQ0cPEi9scSQ+y35AUxKUt/IS2tlpERA1QWO9G+hm+KtOsj/K+qOz/PDkfku+JoPfYN7fUAI91W2
 NSA==
 Container & Folder: myblob/input

Data File Name: employee.txt
Application ID: 4e2e5043-f67b-4bd7-849a-dad4fe1d53ab
Secret Key: Mykdappkey: 3wzRDpdnn3HGZN0vOIGKQ3zIHzbPGFFfVGLInNHZviE=
Tenant ID: cd28ef90-210d-4c10-b290-de42a4d6adfe

OUTPUT



The image shows a Jupyter Notebook interface with a code cell containing JSON output. The output is a list of objects representing pipeline components and their status. The components include 'kirkdataframeapp', 'storageLinkedService', 'ds_in', 'ds_out', and 'copyPipeline'. The 'copyPipeline' object has a 'Pipeline run status' of 'Succeeded' and 'Activity run details' showing 'Activity run status: Succeeded', 'Number of bytes read: 58', 'Number of bytes written: 58', and 'Copy duration: 4'.

```
Name: kirkdataframeapp
Id: /subscriptions/d5f65876-4dec-4a2e-afel-b61e713a8612/resourceGroups/rg-kirkeast/providers/Microsoft.DataFactory/factories/Kirkdataframeapp
Location: eastus
Tags: {}

Name: storageLinkedService
Id: /subscriptions/d5f65876-4dec-4a2e-afel-b61e713a8612/resourceGroups/rg-kirkeast/providers/Microsoft.DataFactory/factories/Kirkdataframeapp/linkedservices/storageLinkedService

Name: ds_in
Id: /subscriptions/d5f65876-4dec-4a2e-afel-b61e713a8612/resourceGroups/rg-kirkeast/providers/Microsoft.DataFactory/factories/Kirkdataframeapp/datasets/ds_in

Name: ds_out
Id: /subscriptions/d5f65876-4dec-4a2e-afel-b61e713a8612/resourceGroups/rg-kirkeast/providers/Microsoft.DataFactory/factories/Kirkdataframeapp/datasets/ds_out

Name: copyPipeline
Id: /subscriptions/d5f65876-4dec-4a2e-afel-b61e713a8612/resourceGroups/rg-kirkeast/providers/Microsoft.DataFactory/factories/Kirkdataframeapp/pipelines/copyPipeline

Pipeline run status: Succeeded

Activity run details

Activity run status: Succeeded
Number of bytes read: 58
Number of bytes written: 58
Copy duration: 4
```

I NOW HAVE OUTPUT FOLDER

myblob
Container

Upload

Refresh

Delete container

Container properties

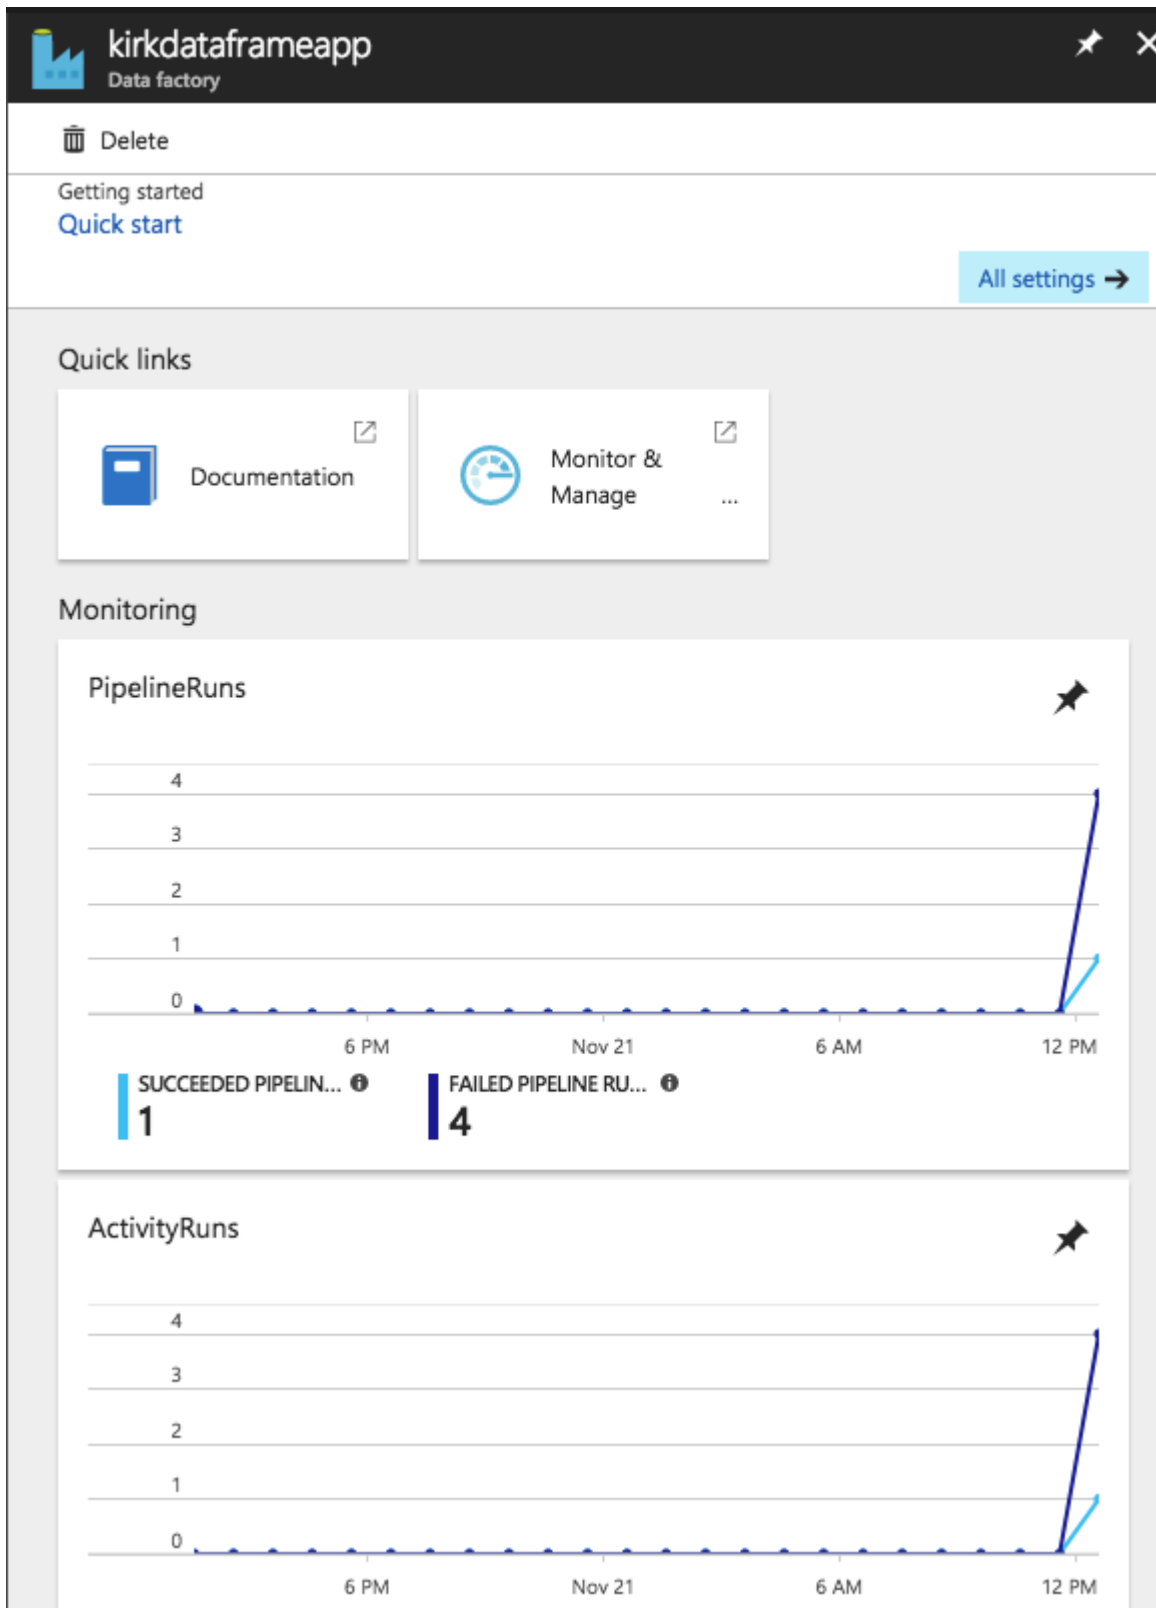
Access policy

Location: myblob

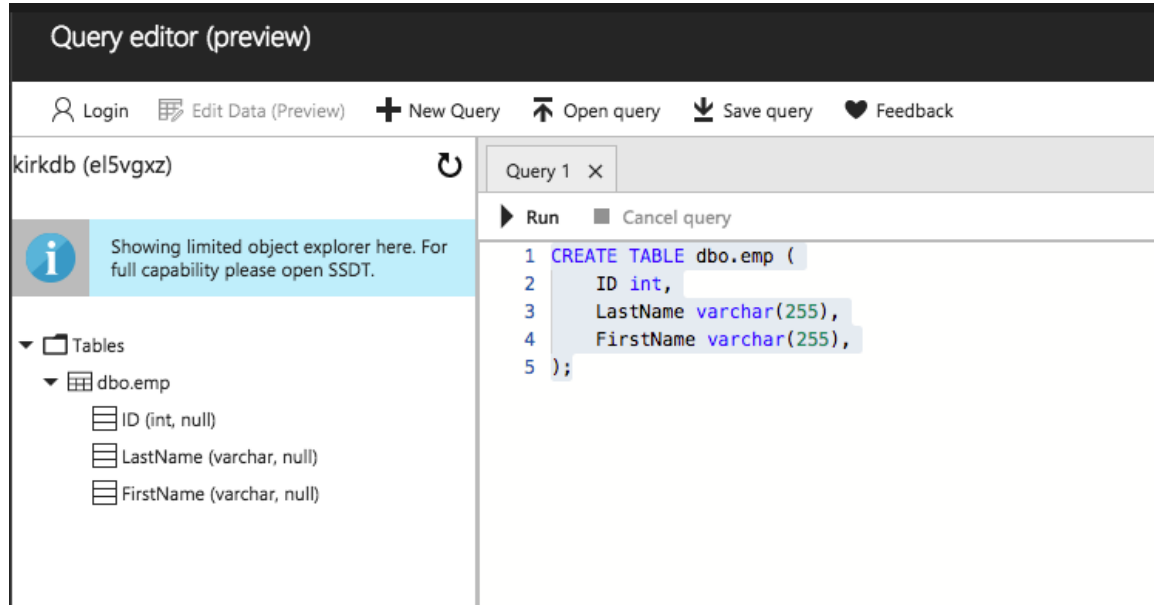
Search blobs by prefix (case-sensitive)

NAME
input
output

SHOWING MY PIPELINE RUNES



CREATED DB



Problem 06. Remove all resource groups created in this assignment using Azure CLI. Please use CLI to show that you have no resources left. (5%)

LIST RESOURCE GROUPS AND DELETE

```
kirks-mbp:week6 el5vgxz$ az group list
[
  {
    "id": "/subscriptions/d5f65876-4dec-4a2e-afe1-b61e713a8612/resourceGroups/cloud-shell-storage-eastus",
    "location": "eastus",
    "managedBy": null,
    "name": "cloud-shell-storage-eastus",
    "properties": {
      "provisioningState": "Succeeded"
    },
    "tags": null
  },
  {
    "id": "/subscriptions/d5f65876-4dec-4a2e-afe1-b61e713a8612/resourceGroups/rg-kirkeast",
    "location": "eastus",
    "managedBy": null,
    "name": "rg-kirkeast",
    "properties": {
      "provisioningState": "Succeeded"
    },
    "tags": null
  }
]
kirks-mbp:week6 el5vgxz$
kirks-mbp:week6 el5vgxz$ az group delete --name rg-kirkeast
Are you sure you want to perform this operation? (y/n): y
kirks-mbp:week6 el5vgxz$ az group delete --name cloud-shell-storage-eastus
```

```
Are you sure you want to perform this operation? (y/n): y  
kirks-mbp:week6 el5vgxz$
```

SUBMISSION INSTRUCTIONS:

Your main submission should be a MS Word or PDF document containing descriptions of your action while configuring Azure services. **If your MS Word document is larger than 1 MB, save it as a MINIMIZED PDF.** Please be merciful and capture small JPGs. Describe the purpose of every action and the significance of the results. Start with the text of this homework assignment as the template. Please add any other files that you might have used or generated. Please write your solution as if you are writing a tutorial for your colleagues. Please make your text readable. Make sure that your fonts, especially in captured images are not unreadable. Please do not provide ZIP or RAR or any other archives. Canvas cannot open those archives and they turn into a nuisance for us.